
***Stanford Shading Group
presentation***

**Hardware Bump-mapping
Choices and Concepts**

May 8, 2000



Mark J. Kilgard

Graphics Software Engineer

NVIDIA Corporation

Coordinate spaces for lighting



Lighting math is rotationally invariant

- Pick the most efficient coordinate space for lighting
- Possible choices
 - eye space (OpenGL per-vertex lighting)
 - world space
 - object space
 - tangent space
- Not screen space since that involves a projection

Evaluating Lighting Coordinate Systems



Eye space

- Advantages
 - eye-space position = view vector
 - nice if hardware floating point dot products are cheap
- Disadvantages
 - lights must be repositioned per-view
 - object-space coordinates must be transformed

Evaluating Lighting Coordinate Systems



World space

- Advantages
 - stationary lights do not need to be repositioned per-view
 - world-space cube maps can encode distant specular illumination
- Disadvantages
 - view vectors harder to compute
 - everything has to be transformed into world space

Evaluating Lighting Coordinate Systems



Object space

- Advantages
 - object space normals need no transformation
- Disadvantages
 - lights and eye must be transformed into object-space
 - view vectors harder to compute

Evaluating Lighting Coordinate Systems



Tangent space

- Advantages
 - normal is always straight up, i.e. $(0,0,1)$
 - makes it easy to perturb for bump mapping!
 - tangent readily available for anisotropic lighting
- Disadvantages
 - lights and eye must be transformed into object-space
 - view vectors harder to compute

Coordinate Systems for Bump Mapping



Common choices

- tangent space
- object space
- but any could work
 - but we want an “easy” space for bump mapping

Tangent Space Bump Mapping



Easy to perturb the normal

- Last column of perturbation rotation matrix “is” the perturbed normal
- By making assumptions, bump map can be “de-coupled” from the object geometry
 - square patch assumption [Peercy 97]
 - saves texture memory since bump maps can be used on different geometry
 - acceptable for morphing geometry

Object Space Bump Mapping



Directly encodes object-space normals

- encode normals directly in texture map
 - but bump map texture is tied to object geometry
 - morphing objects are difficult
 - no need for square patch assumption (though still often assumed)

Encoding Bump Maps



Common choices

- normal maps
 - encode direction vectors (in whatever space)
 - no re-normalization required
- offset maps
 - encode orthogonal offsets from unperturbed normal
 - requires re-normalization
- height fields
 - must approximate derivatives during rendering

Normal maps

Encodes normal vectors directly

- filtering may require re-normalization
 - but not as important
- just perform per-fragment lighting directly using this normal
- diffuse filtering is tractable
- arbitrary normals possible
- scaling issues

Offset maps

Bias unperturbed vector and re-normalize

- cannot encode arbitrary normals
- just Dx, Dy need encoding
 - but large range is needed and filtering is an issue
- requires re-normalization
 - often coupled with cube maps for re-normalization
 - supports so-called “bump environment mapping”
 - but this requires rotations into cube-map space

Height fields

Encodes height of displacement directly

- requires on-the-fly approximation of derivatives
 - usually done with finite differencing or “embossing”
- embossing works poorly when light is straight-on
- very prone to filtering artifacts

Per-fragment Lighting Model



What per-fragment computation

- Usually conventional diffuse/specular/ambient model
- Lambertian diffuse ($L \cdot N$)
 - straightforward
- Specular choices
 - Phong specular ($R \cdot L$)
 - Blinn specular ($H \cdot N$)

Phong Specular Bump Mapping



Often used with offset maps

- Compute un-normalized reflection vector
 - [Voorhies 94]
- Encode specular solution in a cube map
 - rotate perturbed reflection vector into world space
 - cube map provides normalization and exponentiation
 - so-called “free” bump environment mapping
 - do not ignore cost of generating and storing cube map!

Blinn Specular Bump Mapping



Direct approach

- Originally described for bump mapping [Blinn 78]
- Use perturbed normal directly
 - no reflection vector computation
- Requires per-fragment signed dot products
- Evaluating the exponential
 - successive squaring or a look-up table
 - both lead to banding for fixed-point math

Lighting Computation Efficiency



Two approaches

- direct per-fragment lighting computation
 - GPU bump mapping [Kilgard 00]
- pre-computing lighting
 - quantized normal maps [Tarini et.al. 00, Miller et.al. 98]
 - cube maps [Ernst et.al. 95]
 - paletted textures